

Automatic GFL/GFM Reference-Mode Switching in a 1-SG / 4-IBR IEEE-14 System: Model, Small-Signal Spectrum, Switched Time-Domain Formulation, and a Five-Policy Comparison on One Trajectory (250.000 s result set)

Power-flow / stability toolchain — generated report

Abstract

A modified IEEE-14 network is driven through a six-period chronology in which a single synchronous generator (SG) is tripped, the load is stepped, a shunt fault is applied and cleared, a tie line is opened, and the SG is finally reclosed. Four inverter-based resources (IBRs) switch *automatically* between grid-following (GFL) and grid-forming (GFM) reference behaviour under an amplitude–frequency severity index with hysteresis and dwell. This report documents, from the production code, (i) the severity index and its two component sub-indices J_V, J_f , together with the four reference-only diagnostic indices published alongside them; (ii) the sixth-order SG model with an order-by-order derivation and the singular limit that freezes one flux state; (iii) the sixteen-state dual-mode IBR superset with its active projections and current-continuity-gated transfer; (iv) the switched implicit-trapezoidal time-domain formulation, including a proof of the block structure of the stacked right-hand side and the exact mechanism by which an integrated state update is replaced by a discrete transfer at a mode change; (v) the solved power flow, the small-signal spectrum, and the accepted transient response of the 250.000 s run; and (vi) a comparison of *five* control policies — all-GFL locked, one, two and four pinned grid-forming units, and the delivered adaptive supervisor — on the same network, dispatch, event schedule and solver. All equations are those the run executes; all plotted curves are raw accepted samples of that run. No claim of operational readiness is made: convergence of the simulation is not evidence of grid-code compliance.

Contents

1	Introduction	2
2	Scope, sources and reproduction	4
2.1	Nomenclature	5
3	Disturbance chronology	6
4	Automatic GFL/GFM selection: the severity index	7
4.1	Decision evidence: every index on one time axis	8

5	Synchronous generator: sixth-order model and its proof	14
5.1	Order-by-order derivation	14
5.2	Algebraic stator interface	15
5.3	The singular limit that freezes E'_d	15
5.4	Breaker-open coast	15
6	Inverter-based resource: the dual-mode 16-state model	16
6.1	Reduced command-delay states	16
6.2	Shared power stage (states 1–3)	17
6.3	Non-ideal DC source	17
6.4	GFL (states 4–9)	19
6.5	GFM (states 10–16): voltage source, VSG	20
6.6	Current-continuity-gated transfer	20
7	Switched time-domain formulation	21
7.1	Why the device ODEs are not the same	21
7.2	Proof of the block structure of (47)	21
7.3	One coupled DAE, one implicit trapezoidal step	24
7.4	How the step is solved: the damped Newton system	25
7.5	Frozen coordinates: the Π -projection	25
7.6	The switch: from integration to assignment	26
8	Power-flow results	26
9	Small-signal spectrum	27
10	Time-Domain Simulation	29
11	Control-policy comparison	36

1 Introduction

Displacing synchronous machines with converter-interfaced generation removes rotating inertia from the system, and the consequences for frequency stability are by now well characterised: the review in [12] surveys the mechanisms and the control responses, [13] proposes a quantitative inertia-security assessment for exactly this regime, and [14] is one of many synthetic inertia schemes that answer the shortfall from the converter side. What that literature makes clear is that the shortfall is not only a matter of how much inertia is emulated, but of *which converters are willing to form a voltage* when the last machine leaves.

Grid-forming (GFM) control answers that question by making the converter a voltage source with its own angle, the mode whose autonomous operation was characterised early for inverter-based microgrids [9], and grid-following (GFL) control declines it by synchronising to a measured angle through a phase-locked loop. Neither is free. GFM requires active and reactive headroom to be reserved, which is why [15] treats the control mode as a dispatchable grid service with a steady-state cost rather than as a fixed property of a device, and why the capability envelope of a grid-forming inverter has needed explicit characterisation [25]. Under a fault, the two modes behave differently once the current limit binds [24], which is where the comparison stops being a matter of taste.

That cost is what motivates *switching* rather than choosing. Designs that let one converter present either behaviour have been proposed with minimal structural change [16], adapted continuously to the short-circuit ratio [17], made bumpless at the transition [18], or unified into a single grid-following framework that also survives weak grids [19]. In parallel, the transient behaviour of systems that mix machines with both converter types has been analysed directly: [20] for inverter-based generation against a machine, [21] for the damping interaction in a hybrid GFL-GFM pair, [23] for a parallel SG-GFL-GFM system, and—closest to the mechanism this report ends up measuring—[22], which studies what the *switching of* droop-controlled grid-forming converters does to the transient stability of the grid-following converters around them.

What is not settled by that work. Those studies establish that switching helps and that mixed systems have coupled transient behaviour. They do not, between them, answer the question a system operator actually faces on a multi-machine network: given a set of converters that *could* form a voltage, how many should, which ones, and by what route. Two gaps in particular motivate this report. First, a configuration can be admissible in the small-signal sense and still be unreachable from the present state, so a selection rule built on linear screening alone is not sufficient. Second, the number of grid-forming units is usually treated as monotone in benefit, and on the system studied here it is not.

Contributions. This report is a numerical study of one modified IEEE 14-bus system with a single synchronous machine and four converters, driven through a six-period disturbance chronology, and it contributes the following.

1. A complete transcription, from the executing code, of the switched composite model: the severity index that decides the mode (Section 4), the sixth-order machine with its singular-limit flux reduction (Section 5), the sixteen-state dual-mode converter with its current-continuity-gated transfer (Section 6), and the implicit-trapezoidal switched formulation (Section 7).
2. A proof, from stated hypotheses rather than from the implementation, that the composite right-hand side is the direct sum of per-device fields, with a converse showing the hypotheses are necessary and a counterexample drawn from a quantity this system actually computes (Section 7.2).
3. An exhaustive small-signal screening of every candidate grid-forming configuration in the islanded band, which finds only 7 of 15 admissible and the best margin at *two* units rather than four (Sections 9 and 11).
4. A five-policy comparison on one trajectory — all-GFL locked, one, two and four pinned grid-forming units, and the adaptive supervisor — which shows that promoting at least one converter is a necessary condition for a post-trip trajectory to exist at all, that the four-unit configuration fails when pinned from the trip although it passes the linear screen, and that the adaptive arm reaches that same configuration and survives by staging the transitions through per-step certificates (Section 11).

What this report does not claim.

- **No operational readiness.** Convergence of a simulation is not evidence of grid-code compliance, protection coordination or hardware realisability. Transient extrema in what follows are honest excursions of the model under a severe chronology, not operating values or ratings.

- **No selector-against-selector comparison.** The comparison is switching against not switching on *this* resource mix. No claim is made about other selectors, or about systems containing further voltage-forming sources.
- **No multi-machine result.** The repository carries per-machine dynamic data for the reference bus only and the hybrid path refuses several synchronous machines at four fail-closed points, so no numbers for that case appear here and none are assumed in their place.
- **Only S , J_V and J_f decide anything.** J_R , J_P , J_{lock} and J_{SCR} are computed after the trajectory is complete, enter no gate, are classified ASSUMED_DIAGNOSTIC, and must not support a readiness claim. No aggregate is formed from them by design.
- **One linearisation per point.** The spectrum of Section 9 and every damping margin in the admissibility table are single operating-point analyses and do not bound the switched system. Section 11 exhibits a configuration that passes the linear screen and is still not reachable, which is a limit of the linear tool itself.
- **DC-link dynamics are at the resolution limit of the method.** The DC source is non-ideal (Section 6.3) and its mode is faster than the nominal phasor step, so it is resolved only where the adaptive controller reduces the step. A firmer DC source is stiffer still and belongs to an EMT study rather than to this one.
- **Cross-arm numerical quality is not compared.** Residuals and rejected-step counts are not comparable between trajectories of unequal length, and an arm that ends early is never extended to match the others.

2 Scope, sources and reproduction

The study system is the IEEE-14 bus network with one modelled SG at the reference bus and four modelled IBRs at buses 2, 3, 6 and 8. The whole system is assembled as a single all-KCL composite differential–algebraic system

$$g(x, V) = Y_{\text{net}} V - I_{\text{bus}}(x, V) = 0, \quad (1)$$

where Y_{net} is the network admittance, V the stacked complex bus voltages, and I_{bus} the total device injection. The SG and every IBR contribute their own differential state block x and their own injection into (1); no device is delegated to an external solver, model or reference program. Audited linear primitives (`\`, `lu`, `qr`, `eig`) are the only numerical libraries used.

Equation source. Every model equation in Sections 4–7 is transcribed verbatim from the production code that the chronology driver executes. The network, dispatch, event schedule and supervisor gains are fixed by the driver itself, so the whole result set is reproduced by

```
pf_init_paths; run_ieee14_gfm_lock_comparison % t_end=250, dt=0.05
pf_init_paths; generate_ieee14_switch_evidence
```

The first command produces the five arms; the second draws every figure and writes every table and macro this report inputs. The comparison runner builds *one* base option set, adds the opt-in post-processing overlay described next, and then merges a declared per-arm override; it aborts if the two option structures differ in any field that the arm did not declare.

Figure source and no-synthetic-signal contract. The transient figures (Sections 4.1 and 10) are rendered from the accepted trajectory output/diagnostics/ieee14_gfm_lock_

compare/adaptive_250s.mat. Every plotted sample is a raw accepted value. Nothing is smoothed, filtered, decimated, clipped, offset, interpolated or augmented with synthetic noise. Where a display window is applied it is a horizontal range only — no in-window sample is altered or dropped, and an arm whose trajectory ends early is drawn only as far as it ran.

Why this run and not the delivered one. The six diagnostic indices of Section 4.1 are published by an *opt-in* post-processing overlay (option agsi_reference) that defaults to off, so the previously delivered chronology artefact does not contain them, and one of them cannot be recovered after the fact because the admittance log it needs is built under the same flag. The run above therefore re-executes the delivered option set with that single flag added. The re-run was gated on being **bit-identical** to the delivered artefact before any figure was drawn: $\max |\Delta x| = \max |\Delta y| = \max |\Delta u| = 0$ exactly, identical time grid, sample sides, device modes and online flags, identical event log entry by entry, and the same reclose outcome at $t = 159.344$ s. The overlay is therefore established as inert rather than assumed to be. See defect record AGSI-2026-08-16-01.

Classification. The SG AC/flux equations and the IBR AC/control equations are SOURCE_MAPPED: their structural form is the established one in the converter and machine literature cited where each block is derived [1, 2, 6, 7, 10, 11], and their transcription into this project's state layout is recorded in the project's own equation contract. The fixed 16-state IBR super-set, the GFL↔GFM transfer map, the command-delay reduction of Section 6.1 and the DC-source regulator closure are PROJECT_DERIVED. Machine parameters come from the Waterloo IEEE-14 dynamic-data table [4]; the network, branch data and load schedule come from the MATPOWER IEEE 14-bus case [5]; converter parameters, per-unit bases and the disturbance times are otherwise CASE_DEFINED constants of the case file. The predictor and step-subdivision policy are NUMERICAL_METHOD. The four reference-only indices of Section 4.1 are ASSUMED_DIAGNOSTIC and enter no gate. The centre-of-inertia frequency is a project convention, defined explicitly at (55); no external source is claimed for it. These labels are carried from the code provenance fields and are not re-decided here.

2.1 Nomenclature

Table 1 collects the symbols used throughout, grouped by the subsystem that owns them. Symbols are introduced again where they first appear; this table exists so a reader entering at any section can resolve one without searching backwards.

Table 1: Principal symbols.

Symbol	Meaning
<i>System level</i>	
Y_{net}, V	network admittance matrix; stacked complex bus voltages
$I_{\text{bus}}(x, V)$	total device current injection into the network
$g(x, V)$	all-KCL algebraic residual, $Y_{\text{net}}V - I_{\text{bus}}$
x, n_x	composite differential state and its dimension
S_k, E_b	selection operators: $x_k = S_k x$, $V_b = E_b V$
$\mathcal{A}, \mathcal{F}$	active and frozen state index sets, $P_{\mathcal{A}}$ the active projection

Symbol	Meaning
m_k	discrete mode label of device k
$h, r(z), J$	step size; coupled Newton residual; its Jacobian
$S_{\text{base}}, f_0, \omega_b$	system power base; nominal frequency; base electrical speed
κ	system-to-machine power base ratio
f_{coi}	inertia-weighted centre-of-inertia frequency, (55)
Ω	damping margin of a configuration's own equilibrium
<i>Switching supervisor</i>	
J_V, J_f	normalised voltage- and frequency-deviation sub-indices
$S(t)$	severity index, $\text{sat}_{[0,1]}(\frac{1}{2}J_V + \frac{1}{2}J_f)$
$\Gamma_{\text{on}}, \Gamma_{\text{off}}$	upper and lower comparator thresholds
$T_{d,\text{on}}, T_{d,\text{off}}$	minimum dwell before promotion and before release
$J_R, J_P, J_{\text{lock}}, J_{\text{SCR}}$	reference-only diagnostic indices; enter no gate
<i>Synchronous generator</i>	
δ, ω	rotor angle; per-unit speed deviation
E'_q, E'_d	transient EMFs (E'_d frozen at the singular limit)
E''_q, E''_d	sub-transient EMFs
T_m, E_{fd}	mechanical torque; field voltage
T_e, I_d, I_q	air-gap torque; stator currents in the machine frame
H, D, R_a	inertia constant; damping; armature resistance
X_d, X'_d, X''_d	direct-axis synchronous, transient and sub-transient reactances
c_d, d_d, c_q, d_q	flux-coupling coefficients, (13)
<i>Inverter-based resource</i>	
i_d, i_q, V_{dc}	filter currents and DC-link voltage (shared power stage)
$\delta_{\text{PLL}}, \xi_{\text{PLL}}$	PLL angle and integrator (GFL only)
ξ_P, ξ_Q	outer active- and reactive-power integrators (GFL)
$\delta_{\text{VSG}}, \omega_{\text{VSG}}$	virtual-machine angle and speed (GFM only)
E	voltage-loop reference state of the Q - V loop (GFM)
ξ_{Vd}, ξ_{Vq}	voltage-loop integrators (GFM)
ξ_{Id}, ξ_{Iq}	inner current-loop integrators (both modes)
v_{cd}, v_{cq}	inner-loop commanded voltages
R_f, L_f, C_{dc}, T_{dc}	filter resistance and inductance; DC capacitance; DC restoration time
M, D_v, τ_E, k_Q, k_E	virtual inertia, damping, amplitude time constant and Q - V gains
I_{max}, Π_I	current magnitude limit and the circular limiter
$\mathcal{H}$	conditional anti-windup integrator

3 Disturbance chronology

The scheduled disturbance times are CASE_DEFINED constants of the event structure; the re-close and severity-release instants are outcomes read from the accepted event log. The shunt fault is applied at $t = 85.000$ s and cleared at $t = 85.150$ s (a 0.150 s bus-9 fault, impedance

$Z_f = 0.01 + j0.01$ p.u.); this exact window is taken from the driver and supersedes any wider nominal fault interval.

The six operating periods and the transition instants are drawn in Figure 4, which in this revision is deliberately placed *above the result figures* at the head of Section 10, so that the timeline and the curves can be compared in one view without turning back a page.

4 Automatic GFL/GFM selection: the severity index

Each IBR carries a two-state supervisor that decides whether it should behave as a grid-following current source (GFL) or as a grid-forming voltage source (GFM). The decision is driven by a scalar severity $S(t) \in [0, 1]$ built from two normalised sub-indices: a voltage-deviation index J_V and a frequency-deviation index J_f .

Sub-indices. Let $|V_i(t)|$ be the terminal voltage magnitude of IBR i and V_i^* the healthy per-bus power-flow magnitude at the SG-online operating point (Section 8). Let $f_{\text{coi}}(t)$ be the centre-of-inertia frequency. With bases $\Delta V_{\text{base}} = 0.10$ p.u. and $\Delta f_{\text{base}} = 0.50$ Hz,

$$J_V(t) = \frac{||V_i(t)| - V_i^*|}{\Delta V_{\text{base}}}, \quad J_f(t) = \frac{|f_{\text{coi}}(t) - f_0|}{\Delta f_{\text{base}}}, \quad f_0 = 60 \text{ Hz}. \quad (2)$$

Two-term severity (AGSI). The severity is the saturated equal-weight average

$$S(t) = \text{sat}_{[0,1]} \left(\frac{1}{2} J_V(t) + \frac{1}{2} J_f(t) \right), \quad \text{sat}_{[0,1]}(z) = \min\{1, \max\{0, z\}\}. \quad (3)$$

J_V responds to a local voltage collapse or overshoot; J_f responds to a system-wide frequency excursion such as the one that follows the loss of the SG slack at $t = 20$ s. The equal weighting is a PROJECT_DERIVED choice and is frozen before any result is read.

Hysteresis and dwell. A bare threshold on S would chatter. The supervisor therefore uses an asymmetric two-level comparator with minimum dwell times:

$$\begin{aligned} \text{GFL} \rightarrow \text{GFM} : \quad S &\geq \Gamma_{\text{on}} = 0.65 \text{ held for } T_{d,\text{on}} = 0.10 \text{ s}, \\ \text{GFM} \rightarrow \text{GFL} : \quad S &\leq \Gamma_{\text{off}} = 0.35 \text{ held for } T_{d,\text{off}} = 1.00 \text{ s}. \end{aligned} \quad (4)$$

The gap $\Gamma_{\text{on}} > \Gamma_{\text{off}}$ gives the comparator its hysteresis band; the asymmetric dwell ($T_{d,\text{on}} \ll T_{d,\text{off}}$) makes the supervisor quick to add grid-forming support and deliberately slow to release it.

Reference-loss guard. A separate guard runs before the dwell logic and takes precedence over it: a breaker opening that would leave an energised island with *no* voltage-forming device at all is **refused**. The network must not spend even one step without a voltage reference, or (1) has nothing to fix the reference angle and the system loses its mathematical gauge condition. The guard is evaluated *before* Newton is called even once, and Section 11 shows its effect directly: the arm that locks all four IBRs in GFL is refused at exactly $t = 20.000$ s. The guard is a feasibility gate, not a promotion rule — it does not itself raise any converter to GFM. In the delivered run the supervisor reaches its maximum of 4 grid-forming units by 4 separate staged commitments, not by a single simultaneous promotion at the trip.

4.1 Decision evidence: every index on one time axis

Figure 1 puts every published index on one page against one shared time axis, so the value of each can be read at the second the supervisor actually commanded a switch. The vertical lines come in *two families* that must not be confused. Dotted lines are the scheduled disturbances (20, 50, 85, 110, 145 s and the achieved reclose); dash-dotted lines are the supervisor's *own* commitments (mode augmentation and release). In this run the two families do not coincide at a single instant, which is exactly why a figure marked only with the disturbance times cannot show when the switching happened.

Panel (a) is the decision variable itself with both thresholds drawn. Panels (b) and (c) are the two terms that form it. Panels (d)–(g) are reference indices that **enter no gate**: they are computed purely after the trajectory is complete, are classified ASSUMED_DIAGNOSTIC, and must not be used to support a readiness claim. Two properties are worth stating so the figure is not misread: J_f and J_R are *system* quantities built from the centre-of-inertia frequency, so they are single traces rather than four; and J_{lock} is exactly zero on a grid-forming branch by construction, because a GFM branch has no PLL to lose lock — not because the quantity could not be measured. No aggregate index is formed from these four, by design: an aggregate would sit one step away from being a decision variable.

Switching decision --- IEEE 14-bus, 1 SG + 4 converters


adaptive arm, 250 s horizon | decision terms S , J_V , J_f | J_R , J_P , J_{lock} , J_{SCR} reference-only (ASSUMED_DIAGNOSTIC, entered no gate)

Figure 1: Switching-decision evidence; every panel shares one time axis and one set of event lines. (a) Severity $S = \text{sat}_{[0,1]}(0.5J_V + 0.5J_f)$ with the thresholds $\Gamma_{on} = 0.65$ and $\Gamma_{off} = 0.35$. (b) J_V . (c) J_f (system). (d) J_R (system). (e) J_P . (f) J_{lock} . (g) J_{SCR} . (h) per-device GFL/GFM mode. (i) reference-angle owner. Dotted lines: scheduled disturbances. Dash-dotted lines: supervisor commitments. Panels (a)–(c) decide; panels (d)–(g) are reference-only and enter no gate.

The two windows in which the switching actually happens are enlarged in Figures 2 and 3. In those two figures both line families are fully labelled, so the index values at the instant of each command can be read off directly.

Switching decision --- IEEE 14-bus, 1 SG + 4 converters (island formation)

adaptive arm, 250 s horizon | decision terms S, J_V, J_f | $J_R, J_P, J_{lock}, J_{SCR}$ reference-only (ASSUMED_DIAGNOSTIC, entered no gate)

— IBR2 (bus 2) — IBR3 (bus 3) — IBR6 (bus 6) — IBR8 (bus 8)




Figure 2: Islanding window $t \in [17, 58]$ s, covering the SG trip, the first two mode augmentations, one release and the load step. Panels and line families are identical to Figure 1; only the time range differs. No sample is dropped or altered.

Switching decision --- IEEE 14-bus, 1 SG + 4 converters (reclose and handback)

adaptive arm, 250 s horizon | decision terms S, J_V, J_f | $J_R, J_P, J_{lock}, J_{SCR}$ reference-only (ASSUMED_DIAGNOSTIC, entered no gate)




Figure 3: Reclose and hand-back window $t \in [143, 178]$ s, showing the mode releases, one compensating augmentation, the reclose at $t = 159.344$ s, and the return of reference-angle ownership to the SG. Panels and line families are identical to Figure 1.

5 Synchronous generator: sixth-order model and its proof

The SG is a sixth-order round-rotor machine of the structure IEEE Std 1110 recommends for stability studies of this kind [1], with state

$$x_{\text{sg}} = [\delta, \omega, E'_q, E'_d, E''_q, E''_d]^\top, \quad u_{\text{sg}} = [T_m, E_{fd}]^\top, \quad (5)$$

where δ is the rotor angle, ω the per-unit speed deviation (slip), E'_q, E'_d the transient EMFs, and E''_q, E''_d the subtransient EMFs. The two control inputs T_m (mechanical torque) and E_{fd} (field voltage) are held at the constant values determined by the mixed-network equilibrium solve; they are not re-solved each step.

Attribution of the machine structure. Two lineages are easy to conflate here and are kept apart deliberately. The *structural class* — a round-rotor machine carrying transient and sub-transient flux pairs on both axes — is the one IEEE Std 1110 recommends, and that is what the reference above supports [1]. The *particular algebraic form* in which such a machine is coded in industry tools descends from a separate line of work on machine-model equivalence and is not an IEEE 1110 object; the two are not interchangeable and are not treated as such here. IEEE Std 1110 also does not itself supply the swing equation: for that it defers to the standard textbook treatments [2, 3], which is the form used in (6)–(7). No page number is quoted from the standard, because two byte-distinct copies of it circulate and their pagination differs; a page citation that cannot be reproduced is worse than none.

Where the machine numbers come from. The reactances, time constants and inertia of the machine at the reference bus are CASE_DEFINED values taken from the Waterloo technical report's IEEE-14 dynamic-data table [4], mapped onto the flux parameterisation of (13), with $H = 2.5$ s and $D = 1.0$ p.u. on the 100 MVA system base. The network topology, branch data and load schedule are the MATPOWER distribution of the IEEE 14-bus case, itself converted from the IEEE common data format [5]. Neither source is re-derived here and no value from either is adjusted; Section 8 reports the power flow they produce.

5.1 Order-by-order derivation

The connected-machine equations are, term for term as coded,

$$\dot{\delta} = \omega_0 \omega, \quad (6)$$

$$2H \dot{\omega} = T_m - T_e - D \omega, \quad (7)$$

$$T'_{d0} \dot{E}'_q = E_{fd} + c_d E''_q - d_d E'_q, \quad (8)$$

$$T'_{q0} \dot{E}'_d = c_q E''_d - d_q E'_d, \quad (9)$$

$$T''_{d0} \dot{E}''_q = E'_q - E''_q - (X'_d - X''_d) I_d, \quad (10)$$

$$T''_{q0} \dot{E}''_d = E'_d - E''_d + (X'_q - X''_q) I_q, \quad (11)$$

with $\omega_0 = 2\pi f_0$. The air-gap electrical torque that closes the swing equation (7) is

$$T_e = V_d I_d + V_q I_q + R_a (I_d^2 + I_q^2), \quad (12)$$

and the four flux-coupling coefficients are defined from the reactance ladder as

$$c_d = \frac{X_d - X'_d}{X'_d - X''_d}, \quad d_d = \frac{X_d - X''_d}{X'_d - X''_d}, \quad c_q = \frac{X_q - X'_q}{X'_q - X''_q}, \quad d_q = \frac{X_q - X''_q}{X'_q - X''_q}. \quad (13)$$

What each order is. Equations (6)–(7) are the mechanical (2nd-order) swing: angle and speed. Equations (8)–(9) are the transient (rotor field / q-axis damper) fluxes on the slow time constants T'_{d0}, T'_{q0} . Equations (10)–(11) are the subtransient (damper) fluxes on the fast time constants T''_{d0}, T''_{q0} . The armature-reaction terms $-(X'_d - X''_d)I_d$ and $+(X'_q - X''_q)I_q$ couple the fast fluxes to the stator current.

Proof anchor — the $d - c = 1$ identity. From (13),

$$d_d - c_d = \frac{(X_d - X''_d) - (X_d - X'_d)}{X'_d - X''_d} = \frac{X'_d - X''_d}{X'_d - X''_d} = 1, \quad \text{and identically } d_q - c_q = 1. \quad (14)$$

Consider the open-circuit steady state ($I_d = I_q = 0, \dot{x} = 0$). Then (10) gives $E''_q = E'_q$, and substituting into (8),

$$0 = E_{fd} + c_d E'_q - d_d E'_q = E_{fd} - (d_d - c_d) E'_q = E_{fd} - E'_q \implies E'_q = E_{fd}. \quad (15)$$

So the identity (14) is exactly what makes the open-circuit field EMF equal the applied field voltage — the standard consistency check for this flux parameterisation, and the reason the coefficients are written as ratios (13) rather than as free constants. The same argument on the q-axis gives $E'_d = 0$ at open circuit when E_{fd} has no q-axis analogue.

5.2 Algebraic stator interface

The stator currents are *not* reverse-derived from power; they are solved from the subtransient EMFs behind subtransient reactance. Projecting the complex terminal voltage V onto the machine dq frame (V_d, V_q) and writing $r_d = V_d - E''_d, r_q = V_q - E''_q$,

$$\det = X''_d X''_q + R_a^2, \quad I_d = \frac{-R_a r_d - X''_q r_q}{\det}, \quad I_q = \frac{X''_d r_d - R_a r_q}{\det}, \quad (16)$$

which is the exact inverse of the subtransient stator relations $V_d = E''_d - R_a I_d + X''_q I_q, V_q = E''_q - R_a I_q - X''_d I_d$. The electrical torque (12) uses these same I_d, I_q .

5.3 The singular limit that freezes E'_d

For the machine data used here the q-axis transient open-circuit time constant is set to its singular limit $T'_{q0} = 0$ (audited round-rotor case). Then (9) is no longer differential: it degenerates to the algebraic constraint

$$0 = c_q E''_d - d_q E'_d. \quad (17)$$

For this machine $c_q = 0, d_q = 1$, so $E'_d \equiv 0$. The code detects $T'_{q0} = 0$ exactly and *removes* E'_d (state index 4) from the integrated set, holding it at its algebraic value 0. The SG therefore contributes **five** dynamic states $[\delta, \omega, E'_q, E''_q, E''_d]$, not six. This is a SOURCE_DEFINED singular limit, not a numerical convenience: it is the $T'_{q0} \rightarrow 0$ reduction of the same equations.

5.4 Breaker-open coast

When the SG breaker opens at $t = 20$ s the machine does not vanish from the state vector. Its network injection becomes zero ($I_d = I_q = 0, T_e = 0$) and its fluxes decay on the open-circuit branch of (8)–(11) (the armature-reaction terms drop out because $I_d = I_q = 0$); the swing (6)–(7) coasts on a frozen mechanical torque T_m . The map that reports how many states the SG integrates is context-independent: it returns $[1, 2, 3, 5, 6]$ whether the breaker is closed or open. Consequently the loss of the SG slack does *not* change the SG's own state count; the composite dimension changes only because the IBRs switch mode (Section 6).

6 Inverter-based resource: the dual-mode 16-state model

Each IBR is a single device that can present either of two controllers on a common power stage. The full state is a fixed 16-vector,

$$x_{\text{IBR}} = \left[\underbrace{i_d, i_q, V_{dc}}_{\substack{x_{\text{pl}} (3) \\ \text{always live}}} \mid \underbrace{x_{\text{GFL}} (6)}_{\substack{\text{indices 4-9} \\ \text{live only in GFL}}} \mid \underbrace{x_{\text{GFM}} (7)}_{\substack{\text{indices 10-16} \\ \text{live only in GFM}}} \right]^T \in \mathbb{R}^{16}. \quad (18)$$

The two controllers never run at the same time. Equation (18) is a fixed *superset* of coordinates, not a list of simultaneously live states, and this is the single most important thing to read correctly about the device. A converter is *either* grid-following *or* grid-forming at any instant, never both and never partly each: one controller block is integrated while the other is held frozen at the value it carried when it last surrendered the duty, and the two can exchange the duty only at a switching instant, through the transfer of Section 6.6. There is no blend, no weighted average of the two control laws, and no interval on which both sets of integrators advance. The consequence for any later reading of the state vector is concrete: an index in the frozen block is not a measurement of anything during that interval, so a quantity built from a fixed index across a mode change is not a physical signal.

The shorthand $x_{\text{IBR}} = [x_{\text{GFM}} x_{\text{GFL}}]$ is the same object with the shared plant block x_{pl} made explicit and the two controller blocks kept in their code order. The controller blocks are

$$x_{\text{GFL}} = [\delta_{\text{PLL}}, \xi_{\text{PLL}}, \xi_P, \xi_Q, \xi_{Id}, \xi_{Iq}], \quad (19)$$

$$x_{\text{GFM}} = [\delta_{\text{VSG}}, \omega_{\text{VSG}}, E, \xi_{Vd}, \xi_{Vq}, \xi_{Id}, \xi_{Iq}]. \quad (20)$$

Only a projection of (18) is integrated at any instant. The active-index map returns

$$\mathcal{A}_{\text{GFL}} = \{1:9\}, \quad \mathcal{A}_{\text{GFM}} = \{1:3, 10:16\}, \quad \mathcal{A}_{\text{tripped}} = \emptyset, \quad (21)$$

i.e. 9 active states in GFL, 10 in GFM, none when tripped. The intersection $\mathcal{A}_{\text{GFL}} \cap \mathcal{A}_{\text{GFM}} = \{1, 2, 3\}$ is exactly the shared plant, and the union is all of $\{1:16\}$: the plant is always live, and every controller index belongs to one branch alone. The active count is therefore 9 or 10, never 16 — the superset dimension is never the dynamic dimension. With one SG contributing six stored states, the composite state vector of this study therefore has $6 + 4 \times 16 = 70$ rows, which is exactly the row count of the stored trajectory.

6.1 Reduced command-delay states

The source block diagram carries one first-order actuation lag per axis, $T_d \dot{V}_{\text{del}} = v_{\text{cmd}} - V_{\text{del}}$, whose physical time constant is the digital-converter delay $T_d = 1.5/f_{\text{sw}} \approx 0.3$ ms at $f_{\text{sw}} = 5$ kHz. That is more than two orders of magnitude below the phasor step used here, so by singular perturbation the fast lag collapses onto its slow manifold $V_{\text{del}} = v_{\text{cmd}}$: the two delay states per branch are removed algebraically and the filter-current equations use the commanded voltage v_{cd}, v_{cq} directly. This is a standard model-order reduction that leaves both the equilibrium and the retained-state dynamics unchanged, and it removes a pole near 530 Hz that an RMS phasor network model cannot represent in the first place. The reduction takes the dual-mode superset from 20 to the 16 states of (18) and each standalone branch from 12 to 10. It is PROJECT_DERIVED and owner-set; see defect record TD-2026-08-12-01. Report revisions

before this one printed the pre-reduction 20-state layout, which the executed model no longer has.

Two operators recur below; their structural role in a limited converter is the standard one [8], while the particular directional anti-windup realisation is PROJECT_DERIVED. $\Pi_I(a, b; I_{\max})$ is the *circular* current limiter: if $r = \sqrt{a^2 + b^2} > I_{\max}$ it scales $(a, b) \mapsto (a, b) I_{\max}/r$, else it is the identity; its rejected part is the limiter residual $r_i = (a, b) - \Pi_I(a, b)$. $\mathcal{H}(e, \gamma, r_i, \text{sat})$ is the conditional anti-windup integrator: it returns 0 when the reference is clipped and the error would wind further (sat and $\gamma e r_i > 0$), and e otherwise. The instantaneous injection is $I = (i_d + j i_q) e^{j\theta} / \kappa$, $S = V \bar{I}$, $P_{\text{inv}} = \kappa \Re S$, $Q_{\text{inv}} = \kappa \Im S$, with $\kappa = S_{\text{base}}/M_{\text{base}}$ and $(v_d, v_q) = \Re, \Im (V e^{-j\theta})$.

6.2 Shared power stage (states 1–3)

The L -filter current dynamics and DC-link are common to both modes and are always active. Written one state per line in index order:

$$\frac{L}{\omega_b} \dot{i}_d = v_{cd} - v_d - R i_d + \omega L i_q, \quad d\text{-axis filter current} \quad (22)$$

$$\frac{L}{\omega_b} \dot{i}_q = v_{cq} - v_q - R i_q - \omega L i_d, \quad q\text{-axis filter current} \quad (23)$$

$$\dot{V}_{dc} = \frac{1}{C} \left[\frac{E_{dc} - V_{dc}}{R_{dc}} - \frac{P_{ac}}{V_{dc}} - \frac{\max(0, V_{dc} - V_{dc}^{\max})}{R_{ch}} \right], \quad \text{DC-link voltage} \quad (24)$$

with $R = R_f$, $L = L_f$, $C = C_{dc}$, $\omega_b = 2\pi f_b$, and v_{cd}, v_{cq} the inner-loop commanded voltages defined per branch below. The DC-link row is derived in Section 6.3; $P_{ac} = v_{cd} i_d + v_{cq} i_q$ there is the converter-side power, that is the bus power plus the filter loss, which is what the DC bus actually has to supply. The commanded voltage appears here rather than a delayed command because of the reduction of Section 6.1. *Proof of each row.* Rows (1)–(2) are Kirchhoff's voltage law across the series R_f, L_f filter in the rotating frame; the cross terms $\pm \omega L$ are the frame-rotation EMF, and the rotation rate ω is supplied by whichever controller is active (the PLL in GFL, the VSG swing in GFM) — this is the only place the plant equations feel the mode. Row (3) is the PROJECT_DERIVED DC-source closure: the coded current is $I_{dc} = P_{ac}/V_{dc} + (C/T_{dc})(V_{dc}^0 - V_{dc})$ and $\dot{V}_{dc} = (I_{dc} - P_{ac}/V_{dc})/C$, in which the power feed-forward P_{ac}/V_{dc} cancels exactly, leaving the first-order voltage restoration shown. This cancellation is deliberate: the unpublished DC source cannot move the sourced AC equilibrium. At steady state row (3) gives $V_{dc} = V_{dc}^0$.

6.3 Non-ideal DC source

The source publishes the DC-link energy balance $C \dot{V}_{dc} = I_{dc} - P_{ac}/V_{dc}$ but not the law for I_{dc} , so the law is PROJECT_DERIVED and has to be stated and justified rather than assumed. Any DC-side source — a rectifier, a battery, a photovoltaic string in its voltage-source region, or a DC feeder — is to first order an EMF behind its internal plus cable resistance, which gives

$$I_{dc} = \frac{E_{dc} - V_{dc}}{R_{dc}}, \quad (25)$$

and, with the overvoltage chopper of (29) below, the DC-link row of the shared power stage.

E_{dc} is forced, not chosen. Requiring the dispatched point to be an equilibrium, $\dot{V}_{dc} = 0$ at $V_{dc} = V_{dc}^0$ and $P_{ac} = P_{ac}^0$, leaves exactly one admissible value,

$$E_{dc} = V_{dc}^0 + \frac{R_{dc}P_{ac}^0}{V_{dc}^0}, \quad P_{ac}^0 = \kappa P_{\text{ref}} + R_f(i_{d0}^2 + i_{q0}^2), \quad (26)$$

in closed form, because at the initial point both branches set $i_{d0} = \kappa P_{\text{ref}}/|V_0|$ and $i_{q0} = -\kappa Q_{\text{ref}}/|V_0|$ with $\theta = \angle V_0$, so $v_d = |V_0|$, $v_q = 0$, and the two ωL cross terms of $v_{cd}i_d + v_{cq}i_q$ cancel identically. Two consequences matter. First, the sourced AC initial condition and therefore the power flow of Section 8 are untouched: the measured $\dot{V}_{dc}(0)$ is zero to round-off. Second, because (26) is evaluated from the same $P_{\text{ref}}, Q_{\text{ref}}, V_0, R_f$ in both branches, E_{dc} is identical in GFL and GFM, so a run-time transfer introduces no step in $\dot{V}_{dc}$ and not merely none in V_{dc} .

R_{dc} from a declared regulation, bounded by a reachability requirement. Let ε be the DC-source regulation, the fractional DC voltage rise from the dispatched point to no load at rated converter power P_r , so that $R_{dc} = \varepsilon(V_{dc}^0)^2/P_r$. The equilibrium of (25) against a constant-power draw P satisfies $V_{dc}^2 - E_{dc}V_{dc} + R_{dc}P = 0$, which has a real solution only while

$$E_{dc}^2 \geq 4R_{dc}P_{\text{max}}, \quad P_{\text{max}} = V_{\text{max}}I_{\text{max}}, \quad (27)$$

so requiring margin m in (27) bounds ε from above. At $m = 2$ with $V_{\text{max}} = 1.06$ p.u., $I_{\text{max}} = 1.2$ p.u. and the dispatched P_{ref} of this study the bound is $\varepsilon \lesssim 0.115$. The declared value is $\varepsilon = 0.10$, which sits inside that bound with realised margin 2.15 and coincides with the usual ten per cent regulation of a DC link fed through a boost stage or a feeder. It is *not* chosen to suit the integrator step; see the stiffness statement below.

The link cannot run away, and this is proved rather than assumed. For $V_{dc} > E_{dc}$ every term of the DC-link row is negative whenever $P_{ac} \geq 0$, hence $\dot{V}_{dc} < 0$ and

$$V_{dc}(t) \leq \max(V_{dc}(0), E_{dc}) = E_{dc} = V_{dc}^0 \left(1 + \varepsilon P_{ac}^0 / (V_{dc}^0)^2\right) \quad \text{for all } t. \quad (28)$$

When the network collapses during a fault and $P_{ac} \rightarrow 0$, the link therefore rises only towards E_{dc} ; the internal resistance is itself the limit. The overvoltage chopper is a braking resistor conducting in its linear region,

$$I_{ch} = \frac{\max(0, V_{dc} - V_{dc}^{\text{max}})}{R_{ch}}, \quad (29)$$

continuous and Lipschitz, non-differentiable only on $V_{dc} = V_{dc}^{\text{max}}$ in exactly the way Π_I is, so it needs no separate smoothness argument and is covered by Remark 2. By (28) it is inactive whenever $E_{dc} \leq V_{dc}^{\text{max}}$, which holds here; it is retained because a real converter has one and a different dispatch would engage it, and whether it conducted on a given trajectory is decided afterwards from $\max_t V_{dc}$ against V_{dc}^{max} , which is exact because the conduction condition depends on V_{dc} alone.

The stiffness this exposes is reported, not designed around. Linearising the DC-link row with the chopper off about the dispatched point gives

$$\lambda_{dc} = -\frac{1}{C} \left(\frac{1}{R_{dc}} - \frac{P_{ac}^0}{(V_{dc}^0)^2} \right), \quad (30)$$

whose second term is the negative incremental resistance of a constant-power load, so stability requires $R_{dc} < (V_{dc}^0)^2/P_{ac}^0$ — a condition this configuration satisfies by a wide factor. With $C = C_{dc} = 0.10$ and $R_{dc} = 0.10$ the mode is near -95 s^{-1} , a time constant of about 10 ms, so $|\lambda_{dc}| \Delta t \approx 4.8$ at the nominal phasor step $\Delta t = 0.05 \text{ s}$. The implicit trapezoidal rule is A-stable there, but the mode is under-resolved at the nominal step and the adaptive controller has to reduce Δt where it is excited. ε was deliberately not enlarged to bring $|\lambda_{dc}| \Delta t$ near unity: choosing a physical parameter to suit a step size is reverse-engineering it from numerical convenience. The honest consequence is the reverse of a convenience — a *firmer* DC source, ε of order 0.01 for a battery at the terminals, is stiffer still and belongs to an EMT study rather than to a 50 ms phasor step. That is a limit of this class of simulation, stated here rather than hidden by an idealisation.

What this replaced, and why. Revisions of this report before the present one carried a closure with an exact instantaneous power feed-forward, $I_{dc} = P_{ac}/V_{dc} + (C/T_{dc})(V_{dc}^0 - V_{dc})$, which cancels the AC loading identically and collapses the link to $\dot{V}_{dc} = (V_{dc}^0 - V_{dc})/T_{dc}$. Two measured consequences identify it as an ideal DC source. Its own initial condition $V_{dc}(0) = V_{dc}^0$ is the equilibrium of that scalar equation, so V_{dc} stayed at 1.0000000000000000 p.u. for every accepted sample of every converter over the whole 250 s run, range exactly zero — the coordinate carried no information. And because C cancels, the declared capacitance C_{dc} never entered any residual. Both are removed by (25)–(29), in which C_{dc} appears for the first time in (30) and V_{dc} is a live coordinate driven by P_{ac} .

6.4 GFL (states 4–9)

The auxiliary algebraic quantities are the power errors $e_P = \kappa u_1 - P_{\text{inv}}$, $e_Q = \kappa u_2 - Q_{\text{inv}}$, the raw and clipped current references

$$i_{d,\text{ref}}^{\text{raw}} = k_{pP}e_P + k_{iP}\xi_P, \quad i_{q,\text{ref}}^{\text{raw}} = -(k_{pQ}e_Q + k_{iQ}\xi_Q), \quad [i_{d,\text{ref}}, i_{q,\text{ref}}] = \Pi_I(\cdot; I_{\text{max}}), \quad (31)$$

and the inner voltage commands $v_{cd} = v_d + Ri_d - \omega Li_q + k_{pI}(i_{d,\text{ref}} - i_d) + k_{iI}\xi_{Id}$, $v_{cq} = v_q + Ri_q + \omega Li_d + k_{pI}(i_{q,\text{ref}} - i_q) + k_{iI}\xi_{Iq}$. The six GFL states then evolve, one per line in index order:

$$\dot{\delta}_{\text{PLL}} = k_{p,\text{PLL}}v_q + k_{i,\text{PLL}}\xi_{\text{PLL}}, \quad \text{PLL angle; } \omega = 1 + \dot{\delta}_{\text{PLL}}/\omega_b \quad (32)$$

$$\dot{\xi}_{\text{PLL}} = v_q, \quad \text{PLL integrator} \quad (33)$$

$$\dot{\xi}_P = \mathcal{H}(e_P, k_{iP}, r_{i,d}, \text{sat}), \quad \text{active-power integrator} \quad (34)$$

$$\dot{\xi}_Q = \mathcal{H}(e_Q, -k_{iQ}, r_{i,q}, \text{sat}), \quad \text{reactive-power integrator} \quad (35)$$

$$\dot{\xi}_{Id} = i_{d,\text{ref}} - i_d, \quad d\text{-current integrator} \quad (36)$$

$$\dot{\xi}_{Iq} = i_{q,\text{ref}} - i_q, \quad q\text{-current integrator} \quad (37)$$

Proof of each state. (4)–(5) are the synchronous-reference-frame PLL: it drives the measured q -axis voltage v_q to zero, so at lock $\dot{\xi}_{\text{PLL}} = v_q = 0$ and the frame is aligned with the terminal voltage; the loop frequency is $\omega = 1 + \dot{\delta}_{\text{PLL}}/\omega_b$. (6)–(7) are the outer P/Q PIs: at equilibrium

$\dot{\xi}_P = e_P = 0$ forces $P_{\text{inv}} = \kappa u_1$ and $\dot{\xi}_Q = e_Q = 0$ forces $Q_{\text{inv}} = \kappa u_2$, i.e. the IBR meets its P, Q set-points; the $\mathcal{H}$ wrapper freezes them only while the current reference is on the I_{max} circle. (8)–(9) are the inner current PIs: $\dot{\xi}_{Id} = i_{d,\text{ref}} - i_d = 0$ makes the filter current track its reference (and symmetrically on q). The two command-lag states that the source diagram places between v_{cd}, v_{cq} and the filter are not integrated here; they are on their slow manifold by Section 6.1, so $V_{d,\text{del}} = v_{cd}$ and $V_{q,\text{del}} = v_{cq}$ identically — which is the same steady-state relation the lag would have given, reached without carrying an unresolvable fast pole. The GFL branch is a *current source*: it needs the PLL to know the grid angle, so it cannot by itself define a voltage on a dead network.

6.5 GFM (states 10–16): voltage source, VSG

The grid-forming branch replaces the PLL by a virtual-synchronous-generator swing and a Q – V voltage-forming loop. The voltage references are $v_{d,\text{ref}} = E$, $v_{q,\text{ref}} = 0$; the voltage errors are $e_{vd} = E - v_d$, $e_{vq} = -v_q$; the raw/clipped current references are $i_{d,\text{ref}}^{\text{raw}} = k_{pV}e_{vd} + k_{iV}\xi_{Vd}$, $i_{q,\text{ref}}^{\text{raw}} = k_{pV}e_{vq} + k_{iV}\xi_{Vq}$, $[i_{d,\text{ref}}, i_{q,\text{ref}}] = \Pi_I(\cdot; I_{\text{max}})$; and the inner commands v_{cd}, v_{cq} are as in the GFL branch but with $\omega \rightarrow \omega_{\text{VSG}}$. The seven GFM states evolve, one per line in index order:

$$\dot{\omega}_{\text{VSG}} = \omega_b(\omega_{\text{VSG}} - 1), \quad \text{VSG angle} \quad (38)$$

$$M \dot{\omega}_{\text{VSG}} = \kappa u_1 - P_{\text{inv}} - D_v(\omega_{\text{VSG}} - 1), \quad \text{VSG swing} \quad (39)$$

$$\tau_E \dot{E} = k_Q(\kappa u_2 - Q_{\text{inv}}) - k_E(E - E_{\text{ref}}), \quad Q\text{--}V \text{ voltage state} \quad (40)$$

$$\dot{\xi}_{Vd} = \mathcal{H}(e_{vd}, k_{iV}, r_{i,d}, \text{sat}), \quad d\text{-voltage integrator} \quad (41)$$

$$\dot{\xi}_{Vq} = \mathcal{H}(e_{vq}, k_{iV}, r_{i,q}, \text{sat}), \quad q\text{-voltage integrator} \quad (42)$$

$$\dot{\xi}_{Id} = i_{d,\text{ref}} - i_d, \quad d\text{-current integrator} \quad (43)$$

$$\dot{\xi}_{Iq} = i_{q,\text{ref}} - i_q, \quad q\text{-current integrator} \quad (44)$$

Proof of each state. (10)–(11) are the VSG: the swing gives the converter its *own* angle from active-power balance, $M\dot{\omega}_{\text{VSG}} = \kappa u_1 - P_{\text{inv}} - D_v(\omega_{\text{VSG}} - 1)$, so at equilibrium $\omega_{\text{VSG}} = 1$ and $P_{\text{inv}} = \kappa u_1$; because the angle is generated internally there is *no PLL* and the branch can energise a passive network. (12) is the Q – V loop whose state E is the voltage-loop reference $v_{d,\text{ref}} = E$: at equilibrium $E = E_{\text{ref}} + (k_Q/k_E)(\kappa u_2 - Q_{\text{inv}})$, a droop between reactive set-point and formed voltage. (13)–(14) are the voltage PIs that drive $v_d \rightarrow E$, $v_q \rightarrow 0$ (so the terminal voltage aligns with the VSG angle at magnitude E); (15)–(16) the inner current PIs. As in the GFL branch the two command lags are on their slow manifold and are not integrated (Section 6.1). The anti-windup $\mathcal{H}$ again holds the voltage integrators only while the current reference is saturated. The GFM branch is a *voltage source*: this is why at least one grid-forming converter must exist for the island to have a reference angle at all.

6.6 Current-continuity-gated transfer

When the supervisor changes an IBR's mode, the new-mode controller states are not continued by integration; they are *re-initialised* at the current terminal operating point so that the injected current does not jump. Let the pre-switch state x^- inject $I^- = I(x^-, V)$ with $S = V\overline{I^-}$, $P = \Re S$, $Q = \Im S$. The transfer $x^- \mapsto x^+$ is

$$x_{\text{pl}}^+ = x_{\text{pl}}^- (\text{carry } i_d, i_q, V_{dc}), \quad x_{\text{ctrl}}^+ = \text{equilibrium_initialize}(V, P, Q), \quad (45)$$

with one physical carry-across of frequency: GFL→GFM sets $\omega_{\text{VSG}} := \omega_{\text{PLL}}$, and GFM→GFL sets the PLL integrator $\xi_{\text{PLL}} := \omega_b(\omega_{\text{VSG}} - 1)/k_{i,\text{PLL}}$. The transfer is then *gated*: the post-switch

injection $I^+ = I(x^+, V)$ must satisfy

$$|I^+ - I^-| \leq \varepsilon_{\text{tol}} = 10^{-10}, \quad (46)$$

otherwise the step is rejected fail-closed. Equation (46) is the bumpless-transfer guarantee: the mode may change discontinuously in *parameterisation*, but the current the network sees is continuous.

7 Switched time-domain formulation

The composite system is advanced in time by *one* coupled implicit step that the plain and the hybrid drivers share. Three questions are answered in turn, because they are the three things that make this harder than integrating a single machine: (i) the differential equations are *not* the same from one device to the next, nor even for one device across a mode change (§7.1), and the block structure that lets them still be stacked into one right-hand side is proved in §7.2; (ii) they are nonetheless solved *together*, in a single system per step, by discretising the differential part with the implicit trapezoidal rule and appending the network law (§7.3); and (iii) that system is nonlinear, so each step is closed by a damped Newton solve (§7.4). Two refinements specific to this study — the frozen coordinates (§7.5) and the mode switch that replaces integration by a direct assignment (§7.6) — follow.

7.1 Why the device ODEs are not the same

Each device owns a different differential model, so the right-hand side is heterogeneous by construction:

- the **SG** obeys the sixth-order EMF6 equations of Section 5 (swing pair, transient and sub-transient fluxes), of which five are integrated because E'_d is frozen at the singular limit (active set $[1, 2, 3, 5, 6]$);
- an **IBR in GFL** integrates its shared power stage plus the phase-locked-loop controller block — 9 states, $\mathcal{A}_{\text{GFL}} = \{1:9\}$ of (21);
- an **IBR in GFM** integrates the same power stage plus the virtual-synchronous-machine controller (no PLL) — 10 states, $\mathcal{A}_{\text{GFM}} = \{1:3, 10:16\}$;
- a **tripped** device integrates nothing ($\mathcal{A}_{\text{tripped}} = \emptyset$).

The full state is the concatenation of these blocks, $x = [x_{\text{SG}}; x_{\text{IBR}_1}; \dots; x_{\text{IBR}_4}]$, and the right-hand side is the matching concatenation

$$\dot{x} = f(x, V) = [f_{\text{SG}}(x_{\text{SG}}, V); f_{\text{IBR}_1}(x_{\text{IBR}_1}, V; m_1); \dots], \quad (47)$$

where each block f_{IBR_i} depends on that inverter's current mode $m_i \in \{\text{GFL}, \text{GFM}, \text{tripped}\}$ and is written by the device itself. There is no single ODE that all devices share; the coupling between them is *only* through the terminal voltages V that every block reads.

7.2 Proof of the block structure of (47)

Equation (47) asserts something specific about the *model*, not about how the model is stored: that the composite vector field is the direct sum of per-device vector fields, each depending on

the composite state only through its own block. This subsection states the hypotheses under which that is true, proves it, proves the converse so that the hypotheses are seen to be necessary rather than convenient, and then exhibits the one quantity in this very system that would break it.

Notation. Let there be n_d devices, device k attached to bus $b(k)$, carrying a state $x_k \in \mathbb{R}^{n_k}$, an exogenous input u_k , and a mode label m_k with values in a finite set $\mathcal{M}_k$. Put $n_x = \sum_k n_k$ and introduce the selection operators

$$S_k \in \mathbb{R}^{n_k \times n_x}, \quad x_k = S_k x, \quad E_b \in \mathbb{R}^{1 \times n_b}, \quad V_b = E_b V, \quad (48)$$

writing e_b for the b -th standard basis column of $\mathbb{R}^{n_b}$.

Hypotheses.

(A1) Direct-sum state space. The selection operators satisfy

$$\sum_{k=1}^{n_d} S_k^\top S_k = I_{n_x}, \quad S_j S_k^\top = 0 \quad (j \neq k), \quad S_k S_k^\top = I_{n_k}. \quad (49)$$

This is the sentence “the composite state is the ordered concatenation of the device states” written as an operator identity instead of as index arithmetic: the second equation says the blocks are pairwise disjoint, the first says that together they exhaust the composite state.

(H1) One-port coupling. Device k exchanges energy with the rest of the system only through the current at its own terminal. Its internal dynamics therefore depend on the network only through $V_{b(k)} = E_{b(k)} V$.

(H2) Local constitutive law. There exist maps φ_k and ι_k such that

$$\dot{x}_k = \varphi_k(t, x_k, V_{b(k)}, u_k; m_k), \quad I_k = \iota_k(t, x_k, V_{b(k)}, u_k; m_k). \quad (50)$$

(H3) Exogenous mode. On any interval between two consecutive event instants each m_k is constant, and its value is not a function of x .

Proposition 1 (Block structure). *Assume (A1)–(H3) and fix an interval between two consecutive event instants. Then*

(i) *the composite vector field is*

$$f(x, V) = \sum_{k=1}^{n_d} S_k^\top \varphi_k(t, S_k x, E_{b(k)} V, u_k; m_k), \quad (51)$$

which written blockwise is exactly (47); and

(ii) *at every point at which each φ_k is differentiable in its state argument,*

$$\frac{\partial f}{\partial x} = \sum_{k=1}^{n_d} S_k^\top \frac{\partial \varphi_k}{\partial x_k} S_k = \text{blkdiag} \left(\frac{\partial \varphi_1}{\partial x_1}, \dots, \frac{\partial \varphi_{n_d}}{\partial x_{n_d}} \right). \quad (52)$$

Proof. (i) By the first identity of (49), $x = \sum_k S_k^\top S_k x = \sum_k S_k^\top x_k$. Differentiating in t and substituting (50) for each $\dot{x}_k$ gives (51). Multiplying (51) on the left by S_j and using $S_j S_k^\top = \delta_{jk} I$ isolates block j as $S_j f = \varphi_j(t, S_j x, E_{b(j)} V, u_j; m_j)$, so the blocks of f are the per-device fields in device order, which is (47). Note that this part is purely algebraic: it uses no differentiability of any φ_k .

(ii) By (50) the map φ_k depends on x only through $S_k x$, so wherever the derivative exists the chain rule applied to (51) gives the first equality of (52). For the second, take $j \neq k$ and compute the (j, k) block,

$$S_j \frac{\partial f}{\partial x} S_k^\top = \sum_l (S_j S_l^\top) \frac{\partial \varphi_l}{\partial x_l} (S_l S_k^\top) = 0, \quad (53)$$

because by (49) the first factor vanishes unless $l = j$ and the third vanishes unless $l = k$, and $j \neq k$ makes at least one of them vanish in every term. The diagonal blocks are $S_k (\partial f / \partial x) S_k^\top = \partial \varphi_k / \partial x_k$. $\square$

Corollary 1 (Where the coupling lives). *Under the same hypotheses,*

$$\frac{\partial f}{\partial V} = \sum_k S_k^\top \frac{\partial \varphi_k}{\partial V_{b(k)}} E_{b(k)},$$

so the rows of device k see exactly one bus column; and the algebraic constraint is

$$0 = g(x, V) = Y_{\text{net}} V - \sum_{k=1}^{n_d} e_{b(k)} \iota_k(t, S_k x, E_{b(k)} V, u_k; m_k). \quad (54)$$

Hence all interaction between devices is confined to the algebraic block, and two devices meet in it only when $b(j) = b(k)$.

Remark 1 (Solvability is a separate question). Equations (51) and (52) are statements about the map f of two independent arguments x and V . They do not require that V be determined by x ; whether (54) has a unique solution is a different property, enforced at run time by the conditioning gate of Section 7.4 rather than assumed here.

Remark 2 (Why the composite field is only piecewise smooth). Part (ii) is stated pointwise because the composite field is *not* continuously differentiable everywhere: the circular limiter Π_l of Section 6 is not differentiable on the circle $|l| = I_{\max}$, and the conditional integrator $\mathcal{H}$ is not even continuous. Part (i), being algebraic, is unaffected. The finite-difference Jacobian of Section 7.4 inherits the block structure regardless of smoothness: perturbing a coordinate of x_j changes the output of block j alone, by (51), even when that perturbation carries φ_j across one of its own switching surfaces. This is one reason the step is closed with a difference quotient rather than a hand-derived derivative.

Proposition 2 (Converse). *Let $U \subseteq \mathbb{R}^{n_x}$ be open and convex and suppose $f(\cdot, V)$ is C^1 on U with $\partial f / \partial x$ block diagonal with respect to $\{S_k\}$ throughout U . Then for each k the block $S_k f$ is a function of $S_k x$ alone on U .*

Proof. Fix k and let $x, x' \in U$ satisfy $S_k x = S_k x'$. By (49) the difference $d = x' - x$ lies in $\text{span}\{\text{range } S_j^\top : j \neq k\}$. Convexity puts the segment $x + \tau d$, $\tau \in [0, 1]$, inside U , and along it $\frac{d}{d\tau} S_k f(x + \tau d) = S_k (\partial f / \partial x) d = 0$ because every component of d lies in a block whose $(k, \cdot)$ derivative block vanishes. Hence $S_k f(x') = S_k f(x)$. $\square$

Convexity, or a local statement, is needed rather than mere connectedness: on a connected but non-convex U a vanishing directional derivative gives constancy only along each connected component of the corresponding slice.

Remark 3 (The hypotheses have content: a system-wide signal breaks them). Proposition 2 shows the block structure holds *if and only if* the device dynamics are local, so (H2) cannot be dropped. This system in fact contains a quantity that would violate it. The centre-of-inertia frequency is the inertia-weighted mean

$$f_{\text{coi}}(t) = \frac{\sum_{i \in \mathcal{V}(t)} H_i f_i(t)}{\sum_{i \in \mathcal{V}(t)} H_i}, \quad \mathcal{V}(t) = \{i : i \text{ online with } H_i > 0\}, \quad (55)$$

in which $\mathcal{V}(t)$ comprises the synchronous machine and every grid-forming converter, a grid-following converter having no inertia to contribute; and $f_{\text{SG}} = f_0(1 + \omega)$ and $f_{\text{GFM},i} = f_0(1 + \omega_{\text{VSG},i})$ are *states*. Consider the hypothetical modification — *not* the model of this report — in which one converter's swing acquires a system-wide droop term, $M\dot{\omega}_{\text{VSG},k} = \kappa u_1 - P_k - D_v(\omega_{\text{VSG},k} - 1) - D_g(f_{\text{coi}}(x) - f_0)$. That map is not of the form (50), and $\partial(S_k f)/\partial x_j \neq 0$ for every $j \in \mathcal{V}(t)$, so (52) fails and the Jacobian acquires a dense block row. The actual model escapes this because f_{coi} enters only the severity index of Section 4, which acts on the discrete label m_k and therefore through (H3); and because $\mathcal{V}(t)$ is itself determined by the mode vector, which is a second reason the quantity belongs to the switching layer and not to the vector field.

Remark 4 (The object is a switched system). Formally the model is a triple: a finite family of vector fields $\{f^{(m)}\}_{m \in \mathcal{M}_1 \times \dots \times \mathcal{M}_{n_d}}$, a switching signal $t \mapsto m(t)$ that is piecewise constant by (H3), and a reset map applied at the switching instants (Section 7.6). Equation (47) is accordingly a statement about each interval between consecutive events, not about the trajectory as a whole.

Remark 5 (Two operators, not one). The layout operators S_k of (48) are constant: a device's stored state count does not change with its mode. What changes with the mode is the *active* set, carried by a separate projection $P_{\mathcal{A}}$ (Section 7.5). Writing $\mathcal{A}_k$ for device k 's local active set and $o_k = \sum_{j < k} n_j$ for its offset, the composite active set is the disjoint union

$$\mathcal{A} = \bigsqcup_{k=1}^{n_d} (o_k + \mathcal{A}_k), \quad (56)$$

disjoint precisely *because* the S_k are mutually orthogonal by (49). A tripped device contributes $\mathcal{A}_k = \emptyset$, and a singular-limit coordinate such as the E'_d of Section 5.3 is removed from $\mathcal{A}_k$ by its own device.

Hypotheses (A1) and (H1)–(H3) are properties of the device models set out in Sections 5 and 6, not of any particular implementation of them; the provenance of those models is given in Section 2.

7.3 One coupled DAE, one implicit trapezoidal step

The heterogeneous ODE (47) is coupled to the all-KCL network law (1) to form the composite differential–algebraic system

$$\dot{x} = f(x, V), \quad 0 = g(x, V) = Y_{\text{net}}V - I_{\text{bus}}(x, V). \quad (57)$$

It is advanced from (x_0, V_0) to (x_1, V_1) over a step h by discretising the differential rows with the implicit trapezoidal rule and carrying the algebraic rows exactly. With $f_0 = f(x_0, V_0)$,

$f_1 = f(x_1, V_1)$ and the active set $\mathcal{A} \subseteq \{1, \dots, n_x\}$ (the union of the per-device active sets above), the step solves

$$\underbrace{\left[x_1 - x_0 - \frac{h}{2}(f_0 + f_1) \right]_{\mathcal{A}}}_{\text{active trapezoidal rows}} = 0, \quad \underbrace{g(x_1, V_1)}_{\text{all KCL rows}} = 0. \quad (58)$$

Because f_1 and g are evaluated at the unknown end-of-step state, this is an *implicit* (nonlinear, coupled) system, not an explicit update: the states of every device and the voltages of every bus are unknowns of one and the same system, so the machine and all four inverters are solved simultaneously and are never sub-iterated device by device. For an active state i the first block is exactly the trapezoidal rule $x_{1,i} = x_{0,i} + \frac{h}{2}(f_{0,i} + f_{1,i})$.

7.4 How the step is solved: the damped Newton system

Stacking the unknowns as $z = [x_1|_{\mathcal{A}}; V_1]$, the two blocks of (58) become a single nonlinear residual

$$r(z) = \begin{bmatrix} \left[x_1 - x_0 - \frac{h}{2}(f_0 + f_1) \right]_{\mathcal{A}} \\ g(x_1, V_1) \end{bmatrix} = 0, \quad (59)$$

which is driven to zero by a damped Newton iteration. This is the single Newton owner reused by the equilibrium solve, the trapezoidal step, and the event solve; no device supplies its own solver. One iteration is:

1. form the Jacobian $J = \partial r / \partial z$ by *forward finite differences* (column j is $[r(z + \varepsilon e_j) - r(z)] / \varepsilon$ with $\varepsilon = 3 \times 10^{-6}$); no analytic Jacobian is assembled, so a device may change its equations without hand-derived derivatives, and by Remark 2 the difference quotient stays well-defined and block structured where the limiter and anti-windup surfaces make the field non-differentiable;
2. test conditioning: if $\text{rcond}(J) < \epsilon_{\text{mach}}$ the step is rejected fail-closed rather than solved through a singular system;
3. take the Newton direction $\Delta z = -(J \backslash r)$ with the audited MATLAB backslash — the only linear solve in the loop;
4. damp it with a backtracking line search: try $z + \alpha \Delta z$ at $\alpha = 1$ and halve α (up to 20 times) until $\|r\|_{\infty}$ strictly decreases; if no α decreases the residual the step fails closed at the last accepted iterate.

The iteration is declared converged when $\|r(z)\|_{\infty} < \text{tol} = 10^{-8}$, and at most 50 iterations are taken. Nothing in the loop is tuned to the case: the tolerance, the finite-difference increment, the conditioning gate and the line-search contract are fixed constants of the solver.

7.5 Frozen coordinates: the Π -projection

Every index *not* in $\mathcal{A}$ — the frozen set $\mathcal{F} = \{1, \dots, n_x\} \setminus \mathcal{A}$ — is removed from the Newton system and held exactly:

$$x_{1,i} = x_{0,i} \quad \text{for all } i \in \mathcal{F}. \quad (60)$$

The solver begins from $x_1 \leftarrow x_0$ and never writes a frozen row, and a drift check re-verifies $x_1|_{\mathcal{F}} = x_0|_{\mathcal{F}}$ after the step; any change of a frozen coordinate is a hard error. Two distinct things live in

$\mathcal{F}$: the singular-limit flux E'_d of Section 5.3 (frozen for the whole run), and the entire controller block that a given IBR is *not* currently using (frozen for as long as that mode is inactive). The map Π that selects $\mathcal{A}$ is the same runtime authority (`ts_dynamic_state_indices`) at every one of its call sites, so the figure of $n_x(t)$ below cannot disagree with what the integrator actually solved.

7.6 The switch: from integration to assignment

A mode change is landed *exactly* on a step boundary (the step is subdivided so the event time is a grid point). At that boundary the state is not advanced by (58); it is replaced by the discrete transfer of Section 6.6:

$$x_1 = \mathcal{T}(x_0, V) \quad (\text{a re-initialisation}), \text{ not } x_1 = x_0 + \frac{h}{2}(f_0 + f_1). \quad (61)$$

At the mode boundary the controller coordinates jump to the transferred values $\mathcal{T}(x_0, V)$ chosen to satisfy the current-continuity gate (46), and integration resumes on the *new* active set on the very next step. Two important consequences follow for this study:

- The **SG trip** is *not* such an assignment. The SG keeps the same active set [1, 2, 3, 5, 6]; its trajectory is continuous across $t = 20$ s (the coast of Section 5). Only the right-hand side branch changes (connected stator $\rightarrow$ open circuit).
- The **IBR GFL $\rightarrow$ GFM** changes *are* such assignments: each IBR's active dimension steps $9 \rightarrow 10$ and its controller states are re-initialised by (45). The composite dynamic dimension therefore changes because of the inverters, not the generator. In the delivered run these assignments land at the supervisor's own commitment instants (22.052, 51.383, 53.403, 148.267 s), none of which coincides with a scheduled disturbance time.

8 Power-flow results

Table 2 is the solved SG-online power flow (the reference operating point from which V_i^* in (2) is taken); Table 3 the corresponding branch flows and losses. The reference bus supplies $P_g = 1.3319$ p.u., $Q_g = -0.1957$ p.u.; the four IBR buses (2, 3, 6, 8) inject at their scheduled set-points while every load bus meets its P_L, Q_L .

Table 2: Solved bus power flow at the SG-online operating point (per unit on the system base; angles in degrees).

Bus	Type	$ V $	θ (deg)	P_g	Q_g	P_L	Q_L	Q_{min}	Q_{max}	P_{max}
1	REF	1.0000	0.0000	1.3319	-0.1957	0.0000	0.0000	0.00	0.10	3.32
2	GFL	0.9914	-3.2585	0.3172	0.1767	0.2170	0.1270	-0.40	0.50	1.40
3	GFL	0.9667	-8.6921	0.2884	0.1607	0.9420	0.1900	0.00	0.40	1.00
4	PQ	0.9832	-6.7064	0.0000	0.0000	0.4780	-0.0390	–	–	–
5	PQ	0.9858	-5.5221	0.0000	0.0000	0.0760	0.0160	–	–	–
6	GFL	1.0575	-6.7136	0.4443	0.2476	0.1120	0.0750	-0.06	0.24	1.00
7	PQ	1.0253	-6.9205	0.0000	0.0000	0.0000	0.0000	–	–	–
8	GFL	1.0494	-4.4090	0.2677	0.1491	0.0000	0.0000	-0.06	0.24	1.00
9	PQ	1.0219	-8.6416	0.0000	0.0000	0.2950	0.1660	–	–	–
10	PQ	1.0202	-8.5886	0.0000	0.0000	0.0900	0.0580	–	–	–
11	PQ	1.0347	-7.7782	0.0000	0.0000	0.0350	0.0180	–	–	–
12	PQ	1.0411	-7.6862	0.0000	0.0000	0.0610	0.0160	–	–	–
13	PQ	1.0342	-7.8345	0.0000	0.0000	0.1350	0.0580	–	–	–
14	PQ	1.0087	-9.3374	0.0000	0.0000	0.1490	0.0500	–	–	–

Table 3: Branch power flows and losses at the SG-online operating point.

Line	From	To	P_{from}	Q_{from}	P_{loss}	Q_{loss}
1	1	2	0.910918	-0.153112	0.016392	-0.002303
2	1	5	0.420987	-0.042553	0.009593	-0.008907
3	2	3	0.466697	0.013221	0.010470	0.002115
4	2	4	0.316733	-0.064476	0.006065	-0.014738
5	2	5	0.211276	-0.049810	0.002649	-0.025732
6	3	4	-0.197333	-0.018164	0.002803	-0.005014
7	4	5	-0.448842	0.084724	0.002882	0.009089
8	4	7	0.018419	-0.096315	0.000000	0.001990
9	4	9	0.062955	-0.012297	0.000000	0.002223
10	5	6	0.092297	0.001911	0.000000	0.001920
11	6	11	0.130821	0.059907	0.001758	0.003682
12	6	12	0.086088	0.027130	0.000895	0.001864
13	6	13	0.207667	0.085524	0.002983	0.005875
14	7	8	-0.267650	-0.134127	0.000000	0.015018
15	7	9	0.286069	0.035823	0.000000	0.008698
16	9	10	-0.003258	0.021664	0.000015	0.000039
17	9	14	0.057282	0.023368	0.000466	0.000991
18	10	11	-0.093273	-0.036375	0.000790	0.001849
19	12	13	0.024192	0.009266	0.000137	0.000124
20	13	14	0.093739	0.030791	0.001556	0.003168

9 Small-signal spectrum

The modal table below lists the small-signal eigenvalues at a representative post-disturbance configuration (one IBR grid-forming, three grid-following): index, eigenvalue λ , modal frequency f , and the single dominant participating device and state. Conjugate pairs are printed

once. It is the generated table itself, not a re-typeset copy.

No.	Eigenvalue λ (s^{-1})	f (Hz)	Dominant device:state
1	-2075.371441	0	IBR6: i_d
2	-1459.095967	0	IBR8: i_q
3	-1408.957478	0	IBR3: i_d
4	-1288.036461	0	IBR6: i_q
5	-1252.898506	0	IBR8: i_d
6	-952.908217 + j403.267240	64.181975	IBR2: i_q
7	-900.503570	0	IBR2: i_d
8	-249.267967	0	IBR2: ω_{VSG}^{GFM}
9	-159.268727	0	IBR2: E^{GFM}
10	-96.812311	0	IBR2: V_{dc}
11	-92.911157	0	IBR8: V_{dc}
12	-92.693034	0	IBR3: V_{dc}
13	-91.111042	0	IBR6: V_{dc}
14	-13.544228 + j0.055071	8.7648×10^{-3}	IBR2: ξ_{Id}^{GFM}
15	-13.507642	0	IBR2: ξ_{Id}^{GFM}
16	-13.492443	0	IBR8: ξ_{Id}^{GFL}
17	-13.488369	0	IBR6: ξ_{Id}^{GFL}
18	-13.475757	0	IBR3: ξ_{Id}^{GFL}
19	-13.471825	0	IBR8: ξ_{Iq}^{GFL}
20	-13.433124	0	IBR6: ξ_{Iq}^{GFL}
21	-2.623074	0	IBR2: ξ_{Vq}^{GFM}
22	-1.843600	0	IBR8: ξ_p^{GFL}
23	-1.491476	0	IBR8: ξ_Q^{GFL}
24	-1.434890	0	IBR3: ξ_p^{GFL}
25	-1.364330	0	IBR6: ξ_p^{GFL}
26	-1.326920	0	IBR3: ξ_Q^{GFL}
27	-1.185526 + j2.980707	0.474394	IBR2: ξ_{Vd}^{GFM}
28	-0.588912 + j2.171749	0.345645	IBR6: δ_{PLL}^{GFL}
29	-0.566771 + j2.132137	0.339340	IBR3: δ_{PLL}^{GFL}
30	-0.328014 + j0.448860	0.071438	IBR2: ξ_{Vd}^{GFM}

The spectrum is entirely in the left half-plane at this configuration. This is a small-signal statement about one linearisation point; it is not, and must not be read as, a stability certificate for the full switched trajectory.

The DC-link modes are a check on the model, not a free observation. Each converter contributes one DC-link coordinate, and with the non-ideal source of Section 6.3 that coordinate is a live state coupled to the AC states through the converter power P_{ac} . Equation (30) predicts its eigenvalue *before* the spectrum is computed, and with $C_{dc} = R_{dc} = 0.10$ it reduces to the

falsifiable statement

$$\lambda_{dc} = -\frac{1}{C} \left(\frac{1}{R_{dc}} - \frac{P_{ac}}{V_{dc}^2} \right) = -100 + 10 \frac{P_{ac}}{V_{dc}^2} \text{ s}^{-1}, \quad (62)$$

so the four DC rows must be real, must lie a little above -100 s^{-1} , and must be *distinct and ordered by the power each converter carries* — the more heavily loaded the link, the less negative its mode, because the second term of (62) is the negative incremental resistance of a constant-power load. The table satisfies all three: the four rows are separate values with the link voltage as their dominant state, and inverting (62) on them returns per-converter powers in the range 0.3–0.9 p.u., which is the loading the time-domain figures show for the same devices. Nothing was fitted to produce that agreement; the closure was derived first and the spectrum computed afterwards.

Revisions of this report before the present one printed four rows at exactly $\lambda = -10.000000$ with zero modal frequency, unit damping ratio and a participation of exactly 1.000 on a DC-link voltage. That was neither a coincidence nor a numerical artefact but the signature of the earlier ideal closure: it made the DC row exactly decoupled, so $-1/T_{dc}$ was its eigenvalue with algebraic multiplicity four, the unit vector on that coordinate was exactly a left eigenvector, and the participation was therefore exactly one as a matter of algebra. The same idealisation kept V_{dc} identically constant over the whole trajectory. Both are described, with the measurement, at the end of Section 6.3; the closure has been replaced and those four rows are not expected to reappear.

The same analysis over every admissible grid-forming set. The single configuration above is one row of a larger enumeration. With the SG offline, every subset of the four converters that could hold the voltage-forming duty is solved for *its own* equilibrium and linearised on the full-KCL composite system. There are 15 such candidates and only 7 of them are admissible, with the best damping margin at $n = 2$ units ($IBR3 + IBR6$, $\Omega = -0.56814$) rather than at the largest set ($\Omega = -0.48291$). The table and figure are given with the comparison in Section 11, because their purpose there is to show that the admissible sets are *not* monotone in the number of grid-forming units, and that a set can be admissible without being reachable.

10 Time-Domain Simulation

Figure 4 places the six-period timeline directly above every result figure of this section, so that an event instant can be matched against a curve without turning back a page.




Figure 4: Six-period disturbance chronology. Green: undisturbed. Red: a contingency that removes a source or faults a bus. Yellow: a topology change. Blue: a restorative action. Scheduled instants (0, 20, 50, 85.000–85.150, 110, 145 s) are case-defined; the reclose and first severity-release instants are read from the accepted event log.

Table 5 collects the summary scalars of the run that draws every figure in this section. Each is an accepted value of the 250.000 s adaptive arm, read through a generated macro rather than typed into the prose. The voltage and frequency extremes are *reported results* of a deliberately severe event sequence, not acceptance criteria, and no model was adjusted after they were seen.

Table 5: Accepted summary scalars of the 250.000 s run that draws the figures of this section.

Quantity	Value	Note
Last accepted time	250.000 s	full horizon; no truncation
Reclose outcome	SUCCESS	read from the accepted event log
Network voltage range	0.036855–1.523185 p.u.	transient extrema, not ratings
COI frequency range	58.839512–61.138558 Hz	fault and switching window
Terminal COI frequency	60.000000 Hz	at $t = 250.000$ s
Terminal voltage range	0.9667–1.0575 p.u.	at $t = 250.000$ s
Largest accepted step residual	9.9931×10^{-9}	below the 10^{-8} gate at every step
Deepest step subdivision	0 levels	bisection mechanism
Domain-rejected trial steps	0	no gate was relaxed to pass
Controller-rejected steps	137	adaptive stepper
Accepted samples	4994	
SG reclose instant	159.344 s	outcome, not a schedule
Closing mismatch	$\Delta V = 0.008785$ p.u.	within the existing synchronism frame
	$\Delta f = 6.5114 \times 10^{-13}$ pu	
	$\Delta \theta = 5.015017^\circ$	
Mode augmentations (4)	22.052, 51.383,	each passed a transition certificate
	53.403, 148.267 s	
Mode releases (3)	25.310, 146.249,	one device at a time
	152.393 s	
Maximum simultaneous GFM units	4	reached in stages, not at once
Post-reclose reselection status	NO_FEASIBLE_SG_ON_ONE_STEP	a correct fail-closed outcome

Figure 5 reads top to bottom as cause and then effect. Panel (a) is the severity index S of Section 4 with both thresholds drawn, one trace per converter because J_V is each converter's own terminal deviation; it is the only quantity on the page that *causes* anything. Immediately below it are two supervisory strips: each IBR's grid-forming/grid-following mode, and the *reference owner* — the single device that holds the island angle reference at each instant. Panels (b) and (c) are the per-converter active and reactive power over the full horizon 0–250.000 s. Mode and reference ownership are distinct contracts: several IBRs can be grid-forming at once, but only one of them owns the reference. Here the SG owns it up to $t = 20$ s and again after the reclose; in between, IBR₁ (bus 2) is the reference leader.

The point of this result set is in the mode strip: the supervisor does *not* raise all four converters together. It augments in 4 stages, at 22.052, 51.383, 53.403 and 148.267 s, until it holds 4 grid-forming units, then releases 3 times after the reclose. Section 11 shows that *pinning* those same 4 units from the trip onwards destroys the run before the horizon, even though that set passes the linear admissibility check. The order in which the set is reached therefore matters as much as the set itself.




Figure 5: (a) Severity index S per converter with both switching thresholds, $\Gamma_{on} = 0.65$ and $\Gamma_{off} = 0.35$. Second strip: grid-forming or grid-following mode of each IBR. Third strip: the single device holding the island angle reference. (b) Per-IBR active power P_i . (c) Per-IBR reactive power Q_i , both per unit. Dotted lines are scheduled disturbances and dash-dotted lines the supervisor's own commitments; near-coincident instants share one label. Raw accepted samples over the full 250.000 s horizon.

Figure 6 widens the same trajectory into the eight electrical quantities that the device reconstruction publishes. Panel (h) is a whole-network scalar and therefore has one trace by definition; every other panel carries the four converters and, where a synchronous counterpart exists, the SG.




Figure 6: Eight-panel electrical response over the full 250.000 s horizon. (a) active power; (b) reactive power; (c) d -axis current; (d) q -axis current; (e) PCC or virtual-rotor frequency; (f) device-to-PCC angle; (g) PCC voltage; (h) network minimum voltage. The SG is dashed black wherever a synchronous counterpart exists; panel (h) is a system scalar and has a single trace by definition. Vertical lines mark the same instants as Figure 5, unlabelled here because each panel is half a page wide. Raw accepted samples.

11 Control-policy comparison

The same trajectory is run five times with the *same* network, dispatch, event schedule, solver and step controller, changing only the IBR control policy. The runner declares which option fields each arm may differ in and checks that the structure it built differs from the shared base in no other field. All five arms are bit-identical on $[0, 20.000)$ s, before the first event, with $\max |\Delta x| = \max |\Delta y| = \max |\Delta u| = 0$ over the 42 shared samples, so every later difference is an effect of the control policy alone. The figures and tables below are the result; each caption states what it shows and what may not be concluded from it.

Table 6: Outcome of five control policies on the same 250.000 s trajectory. A – means the quantity is meaningless for that arm, not that it is zero.

Control policy	Reached [s]	Reclose [s]	Max GFM	Refused by	Outcome
all IBRs locked GFL	20.000	–	0	noVoltageFormingSource	island left with no voltage-forming source; refused before Newton runs
1 GFM pinned (IBR ₁)	250.000	151.2760	1	–	survives the horizon, but carries the whole imbalance alone
2 GFM pinned (IBR ₂ +IBR ₃)	250.000	153.7212	2	–	survives the horizon; best-damped set in the table
all 4 GFM pinned	25.490	–	4	adaptiveDtMin	fails before the horizon , although this set is linearly admissible
automatic switching (delivered)	250.000	159.3436	4	–	survives the horizon <i>and</i> reaches the four-unit set

comparison_switch_vs_lock.png
 pending figure regeneration.

Figure 7: Switching against not switching, in the window where the two part company. Rows, top to bottom: centre-of-inertia frequency, network minimum voltage, number of grid-forming units, and which device holds the island angle reference. **Switching** (solid) rides through the loss of the synchronous machine and continues to the end of the horizon. **All four units pinned grid-forming** from the outset, with no run-time switching, breaks into a growing oscillation and the trajectory ends at $t = 25.490$ s with `adaptiveDtMin`; the terminating sample is marked, and the curve is not continued past it. **All units locked grid-following** cannot be integrated at all past the trip: the island is left with no voltage-forming source, the per-island admissibility gate refuses the transaction before Newton is called even once, and the trajectory ends at the left-limit sample of the event, exactly $t = 20.000$ s. An arm that ends early is never extended to match the others, and no quantity is reported for an arm over an interval it did not reach.




Figure 8: Five arms overlaid over the full horizon. Rows top to bottom: centre-of-inertia frequency, network minimum voltage, number of grid-forming units, and reference-angle owner. The chronology is named once, on the top row, in a band reserved above the data. This page reads the *settled* behaviour: the three surviving arms separate into three distinct frequency levels, and the staged $1 \rightarrow 2 \rightarrow 4$ progression of the switching arm is visible in the third row. An arm that ends before the horizon stops at its own last accepted sample (solid marker) and is *not* extended, held, or extrapolated to match the others. Every y range is taken from the data actually inside the plotted time window, and the generator aborts if any sample would fall outside it.




Figure 9: The same four rows over the three windows in which the transients are resolvable, every bound derived from the event schedule. Left: island formation, where the pinned four-unit arm breaks into a growing oscillation and terminates. Centre: the fault, with the 0.150 s shunt interval shaded, showing the network minimum voltage collapsing and recovering. Right: reclose and hand-back. Each column is scaled to its own window, which is what the previous single-page version failed to do.

Table 7: Frequency deviation in the 15-second window after each event, for the three arms that reach the horizon. Peak is $\max |f_{COI} - 60|$; settled is the mean of $|f_{COI} - 60|$ over the last 30% of the window. All values in hertz.

Event	Peak			Settled		
	1 GFM	2 GFM	switching	1 GFM	2 GFM	switching
SG trip (20 s)	0.5312	0.3489	0.7587	0.0113	0.0001	0.0040
load +20% (50 s)	1.4478	0.7839	1.0042	1.2215	0.7124	0.3753
fault (85 s)	1.4977	0.8257	1.1386	1.2296	0.7125	0.3753
line 6–13 open (110 s)	1.2325	0.7135	0.3764	1.1426	0.6422	0.3503
reclose (145 s)	1.1736	0.6634	1.1605	0.0203	0.0232	0.0120




Figure 10: (a) Peak frequency deviation in the 15-second window after each event. (b) Settled deviation in the same window. (c) Network minimum voltage. A missing bar means that arm never entered that window; it does not mean zero.

Table 8: Every candidate grid-forming subset for the islanded (SG-off) band, screened before the run. Ω is the damping margin of that configuration's own equilibrium; a more negative value is better damped. The admissible set is *not* monotone in the number of grid-forming units, and the best two-unit configuration is better damped than all four.

grid-forming set	reference	n	Ω	outcome
IBR8	IBR8	1	-0.15544	admissible
IBR6	IBR6	1	-0.25203	admissible
IBR3	IBR3	1	-0.31671	admissible
IBR2	IBR2	1	-0.32801	admissible
IBR2+IBR6	IBR2	2	-0.56172	admissible
IBR3+IBR6	IBR3	2	-0.56814	admissible
IBR2+IBR3	IBR2	2	–	no equilibrium (IBR6 at a limit)
IBR2+IBR8	IBR2	2	–	no equilibrium (IBR6 at a limit)
IBR3+IBR8	IBR3	2	–	no equilibrium (IBR6 at a limit)
IBR6+IBR8	IBR6	2	–	no equilibrium (IBR3 at a limit)
IBR2+IBR3+IBR6	IBR2	3	–	no equilibrium (Newton did not converge)
IBR2+IBR3+IBR8	IBR2	3	–	no equilibrium (Newton did not converge)
IBR2+IBR6+IBR8	IBR2	3	–	no equilibrium (Newton did not converge)
IBR3+IBR6+IBR8	IBR3	3	–	no equilibrium (Newton did not converge)
IBR2+IBR3+IBR6+IBR8	IBR2	4	-0.48291	admissible

Best authenticated configuration. IBR3+IBR6, with $n = 2$ and $\Omega = -0.56814$. The best singleton reaches $\Omega = -0.32801$ and the full set -0.48291, so the optimum is interior: adding grid-forming units past it degrades the margin. This is a small-signal result about one equilibrium and does not by itself certify the transient response; the arm comparison measures that separately.




Figure 11: Damping margin Ω of every *admissible* grid-forming configuration, each at its own measured value, against the number of units (more negative is better damped). The connecting line follows the best configuration at each cardinality, which makes the interior optimum at two units immediate. Labels are separated in point space with leader lines where two margins differ by less than one line of type. The configurations with no equilibrium are *not* plotted: a marker at a position that is not its value cannot be measured off the axis. Their count at each cardinality is printed under the axis and their individual reasons are in Table 8.

References

- [1] IEEE Std 1110-2002, *IEEE Guide for Synchronous Generator Modeling Practices and Applications in Power System Stability Analyses*.
- [2] P. Kundur, *Power System Stability and Control*. New York: McGraw-Hill, 1994.
- [3] P. W. Sauer and M. A. Pai, *Power System Dynamics and Stability*. Upper Saddle River, NJ: Prentice Hall, 1998.
- [4] S. K. M. Kodsí and C. A. Cañizares, “Modeling and simulation of IEEE 14 bus system with FACTS controllers,” Univ. of Waterloo Technical Report 2003-3, 2003. (Machine dynamic data, Table A.2.)
- [5] MATPOWER case case14, IEEE 14-bus test case converted from the IEEE common data format, as distributed with the MATPOWER package. Used here for network topology, branch data and the load schedule.
- [6] A. Yazdani and R. Iravani, *Voltage-Sourced Converters in Power Systems: Modeling, Control, and Applications*. Hoboken, NJ: Wiley/IEEE Press, 2010, ISBN 978-0-470-52156-4.
- [7] R. Teodorescu, M. Liserre, and P. Rodríguez, *Grid Converters for Photovoltaic and Wind Power Systems*. Chichester, UK: Wiley/IEEE Press, 2011, ISBN 978-0-470-05751-3.

- [8] S. Bacha, I. Munteanu, and A. I. Bratcu, *Power Electronic Converters Modeling and Control*. London: Springer, 2014, doi:10.1007/978-1-4471-5478-5.
- [9] N. Pogaku, M. Prodanović, and T. C. Green, “Modeling, analysis and testing of autonomous operation of an inverter-based microgrid,” *IEEE Trans. Power Electron.*, vol. 22, no. 2, pp. 613–625, 2007.
- [10] K. Sakimoto, Y. Miura, and T. Ise, “Virtual synchronous generator without phase locked loop based on current controlled inverter and its parameter design,” *IEEE Trans. Power and Energy*, vol. 135, no. 7, pp. 462–470, 2015.
- [11] National Renewable Energy Laboratory, *REGFM_B1: An Electromagnetic Transient Model for Grid-Forming Inverters*, NREL/TP-5D00-90260, 2024.
- [12] C. He, H. Geng, K. Rajashekara, and A. Chandra, “Analysis and control of frequency stability in low-inertia power systems: a review,” *IEEE/CAA J. Autom. Sinica*, vol. 11, no. 12, pp. 2363–2383, Dec. 2024, doi:10.1109/JAS.2024.125013.
- [13] W. Zhang, Y. Wen, and C. Y. Chung, “Inertia security evaluation and application in low-inertia power systems,” *IEEE Trans. Power Syst.*, vol. 40, no. 2, pp. 1725–1737, Mar. 2025, doi:10.1109/TPWRS.2024.3420786.
- [14] S. Nema, V. Prakash, and H. Pandžić, “Adaptive synthetic inertia control framework for distributed energy resources in low-inertia microgrid,” *IEEE Access*, vol. 10, 2022, doi:10.1109/ACCESS.2022.3177661.
- [15] G. Cui, Z. Chu, and F. Teng, “Control-mode as a grid service in software-defined power grids: GFL vs GFM,” *IEEE Trans. Power Syst.*, vol. 40, no. 1, pp. 314–326, Jan. 2025, doi:10.1109/TPWRS.2024.3404339.
- [16] H. Ding, R. Kar, Z. Miao, and L. Fan, “A novel design for switchable grid-following and grid-forming control,” *IEEE Trans. Sustain. Energy*, vol. 16, no. 2, pp. 1301–1314, Apr. 2025, doi:10.1109/TSTE.2024.3520989.
- [17] P. Liu, X. Xie, and J. Shair, “Adaptive hybrid grid-forming and grid-following control of IBRs with enhanced small-signal stability under varying SCRs,” *IEEE Trans. Power Electron.*, vol. 39, no. 6, pp. 6603–6607, Jun. 2024, doi:10.1109/TPEL.2024.3373659.
- [18] X. Gao, D. Zhou, A. Anvari-Moghaddam, and F. Blaabjerg, “Seamless switching method between grid-following and grid-forming control for renewable energy conversion systems,” *IEEE Trans. Ind. Appl.*, vol. 61, no. 1, pp. 597–606, Jan./Feb. 2025.
- [19] A. Askarian, J. Park, and S. Salapaka, “Enhanced grid-following (E-GFL) inverter: a unified control framework for stiff and weak grids,” *IEEE Trans. Power Electron.*, vol. 39, no. 5, pp. 5089–5107, May 2024, doi:10.1109/TPEL.2024.3350528.
- [20] C. He, X. He, H. Geng, H. Sun, and S. Xu, “Transient stability of low-inertia power systems with inverter-based generation,” *IEEE Trans. Energy Convers.*, vol. 37, no. 4, pp. 2903–2912, Dec. 2022, doi:10.1109/TEC.2022.3185623.
- [21] Q. Qu, X. Xiang, K. Xin, Y. Liu, W. Li, and X. He, “Transient stability analysis of hybrid GFL-GFM system considering various damping effects,” *IEEE Trans. Ind. Electron.*, vol. 72, no. 12, pp. 13334–13347, Dec. 2025, doi:10.1109/TIE.2025.3581260.

- [22] B. Li, S. Yang, P. Cheng, and Z. Hao, “Transient stability of GFL converters subjected to switching of droop-controlled GFM converters,” *IEEE Trans. Energy Convers.*, early access, doi:10.1109/TEC.2026.3676254. (Author’s accepted version; content may change before final publication.)
- [23] L. Chen, S. Zheng, R. Tao, Z. Liu, Y. Li, Y. Li, Y. Jiang, and H. Chen, “Transient stability analysis approach of the SG-GFL-GFM parallel system based on power-angle interaction mechanism,” *Protection and Control of Modern Power Systems*, early access, doi:10.23919/PCMP.2025.000197.
- [24] Z. M. Almesri, H. A. Hussain, and R. M. Kamel, “Low voltage ride-through capability of grid-connected PV systems: a comparative study of grid-following and grid-forming converters under current limits,” *IEEE Access*, vol. 13, pp. 105590–105607, 2025, doi:10.1109/ACCESS.2025.3580643.
- [25] M. Nurunnabi, S. Li, and H. S. Das, “Advancing grid-forming inverter technology: comprehensive PQ capability and performance analysis,” *IEEE Access*, vol. 13, pp. 73391–73407, 2025, doi:10.1109/ACCESS.2025.3561788.